

Programming Assignment 3 Foundations of Computational Math 1 Fall 2025

The solutions are due on Canvas by 11:59 PM on Sunday, October 26, 2025

Mandatory Meetings

You will meet with the TA on either Wednesday October 22 or Thursday October 23 for a progress check. You must bring your in-progress code. Be prepared to discuss your work so far and bring any questions/clarifications regarding the assignment. Note that the assignment does not need to be complete. You need to demonstrate some progress and understanding of the assignment/your code. Attending the meeting and demonstrating your progress and understanding is worth 5 points of the assignment grade. Failure to sign-up and attend or coming unprepared to the check-in with no progress will result in a loss of those points. The meeting time sign-up will be at the FCM recitation on Monday October 20. If you are unable to attend recitation, it is your responsibility to contact the TA to select a remaining available meeting time.

General Task

This programming assignment implements and empirically evaluates the transformation of a full-rank matrix $A \in \mathbb{R}^{n \times n}$ to an upper triangular form, i.e.,

$$T_L^{-1}A \rightarrow U$$

where $U \in \mathbb{R}^{n \times n}$ and nonsingular. You must produce code to use Gauss transformations and row permutations to achieve this transformation

$$P_r A = LU$$

where $P_r \in \mathbb{R}^{n \times n}$ is a permutation matrix based on a series of interchanges of pairs of rows, $L \in \mathbb{R}^{n \times n}$ is a unit lower triangular matrix, and $U \in \mathbb{R}^{n \times n}$ is a nonsingular upper triangular matrix.

It is not required but you are encouraged to then add the capability to your code to apply complete pivoting, i.e., computing

$$P_r A P_c = LU$$

where P_c is a permutation based on a series of interchanges of pairs of columns. Extra credit will be awarded for a successful implementation and evaluation of complete pivoting.

This assignment empirically evaluates the predictions you make about the structure of the factors and the general performance of the method across a large set of randomly generated problems and selected specific problems.

The Codes

1. Implement a code that computes the LU factorization of a matrix A without pivoting first, then add partial pivoting (rows), and finally add complete pivoting if you wish. You should use the immediate update form of algorithm based on rank-one updates. The final code should be capable of performing the various tasks, based on a user specified parameters. Your code should also detect situations where the factorization may not proceed and exit gracefully. Clearly, this is the case when the candidate pivot set contains no acceptable value. In exact arithmetic this means they are all 0. Your code should also allow the detection of a set of candidate pivots that are all “too small” and warn the user. For no pivoting this would be the $(1,1)$ element of the current active part of the matrix is below some threshold; for partial pivoting this would be all elements in the first column of the current active part of the matrix are below some threshold; for complete pivoting this would be all elements in the entire active part of the matrix are below some threshold.

Your factorization routine, which must be a separate routine from the driver/tester routine, should accept the matrix A stored in a simple 2-dimensional array-like data structure and other relevant parameters such as n and a flag indicating what form of the factorization should be attempted, e.g., no pivoting, partial pivoting, or complete pivoting. The routine should return the matrices L and U stored within the array that contained A on input, i.e., you should implement the in-place algorithm strategy described in the class notes and lectures. You should also keep a copy of A in an additional data structure for correctness checking. The 1 values on the diagonal of L , the 0 values in the upper triangular portion of L and the 0 values in the lower triangular portion of U should not be stored at any point in the code. The routine should also return the permutation matrix P_r when partial pivoting is used and P_r and P_c if you implement complete pivoting. These permutation matrices should not be stored as a full matrix within any array. A permutation matrix or equivalently the set of elementary permutations that are its factors, can be represented by at most n integers.

2. You will also need various test routines designed to evaluate and validate the correctness of the code and accomplish the tasks described below. You may code in Matlab, Python (as long as you write basic code that exposes data structure details and vector or scalar loop constructs) or any compiled and typed language you wish although C, C++, Julia, and Fortran are preferred. In all cases, however, you may not use standard libraries such as LAPACK or built-in matrix routines for pieces of your routines implementing the computations described above. (They may be used for the validation codes.)
3. For both the factorization, triangular solutions and testing codes you should make use of your solutions to programming assignment 1, the posted solutions and other posted example codes. If you were unable to get any of them working reliably then you are free to use or convert my solutions. The testing drivers from my posted solutions should

also be adapted or used to inspire your test driver's design for this assignment. A driver to test LU factorization has also been posted to assist your solution development.

Routines to support the evaluation of the factorization include but are not limited to:

1. Implement a routine that accepts as input the one or two 1-dimensional arrays specifying P_r and P_c and applies them to a matrix stored in a 2-dimensional array and returns the result in a separate 2-dimensional array, i.e., $A_2 \leftarrow P_r A_1 P_c$.
2. Implement, or use the posted solution, a routine that accepts as input a vector b stored in a simple 1-dimensional array, input the 2-dimensional array containing the information specifying the L and U matrices and the 1 dimensional arrays specifying P_r and P_c , if appropriate, from the factorization of a matrix A and returns the solution x of the system $Ax = b$. Any application of the permutations should use the permutation routine listed earlier. The forward and backward solves $Ly = b$ and $Ux = y$ required in this routine must also work only with the information specified in the 2-dimensional array and not expanded versions of L and U .
3. Implement, or use your code from your Program 1 solutions or the posted solution, a matrix multiplication routine that accepts as input the 2-dimensional array containing the information specifying the L and U matrices and returns in a separate 2-dimensional array the product $M = LU$ or the product $M = |L||U|$ based on user selection indicated by an input flag. (This second function requires an addition to your routines from Program 1 if you use them.) Note that these matrix multiplications must use the information specifying L and U within the data structure. It **must not** expand them into separate arrays that include the 1 and 0 values that are not stored in the input array. The notation $|B|$ defines a matrix with elements equal to the absolute values of corresponding elements of B , i.e., $e_i^T |B| e_j = |e_i^T B e_j|$.
4. The test routines should include the computation of $\|W\|$ where W is a matrix. The norms used should be finite in computation, i.e., $\|W\|_1$, $\|W\|_\infty$, or $\|W\|_F$. You may use library routines to compute $\|W\|_2$ if you wish but you are definitely not expected to generate the code for the 2-norm.

Library Codes

You may use libraries and external routines in your test drivers to generate solutions for comparisons, to generate histograms, graphs and any other useful summary display mechanisms. Make sure when using library routines as part of your empirical analysis, you carefully check, e.g., P_r , P_c , L and U generated by your routines and those generated by the library, e.g., MATLAB. These factors are not unique given that pivoting choices are not unique in general. So they may not match. The products LU however should match the matrices $P_r A P_c$ generated by your routine and a library respectively.

Metrics

There are several important metrics to use when assessing the code's correctness. These metrics should be computed in double precision especially if you have run your routines in single precision. (The ability to run in both is not required for this assignment but is easily done if you use a compiled and typed language and is, in general, a good idea to design into your code.)

Some suggestions follow:

1. When comparing matrices use one of the finitely computable ones, $\|M\|_1$, $\|M\|_\infty$ and $\|M\|_F$ which you can easily implement. You may also use a library function from Matlab etc. to compute the matrix 2-norm if you wish.
2. Check the factorization accuracy

$$\frac{\|P_r A P_c - LU\|}{\|A\|}$$

where $\|A\| \geq 1$, i.e., relative error for large A , and $P_c = I$ for partial pivoting and $P_r = P_c = I$ for no pivoting.

3. If the solution is known by design of the problem check

$$\frac{\|x - \tilde{x}\|}{\|x\|}$$

where \tilde{x} is the computed solution and $\|x\| \geq 1$.

4. You should check the accuracy via the residual $b - A\tilde{x}$ and

$$\frac{\|b - A\tilde{x}\|}{\|b\|}$$

assuming $\|b\| \geq 1$ for all attempts to solve a system, i.e., whether or not you know the true solution.

5. If you have access to a standard library you may also use the results of its LU factorization algorithms. However, care must be taken since details of pivoting strategies may yield differences in the factors and permutations. The library routines or your Program 1 solutions are very useful when you generate a system by choosing A and b and need a reliable way of generating x to compare with your computed solution.
6. You should compute the following scalar known as the growth factor

$$\gamma_\epsilon = \frac{\| |L_\epsilon| |U_\epsilon| \|}{\|A\|}$$

where $L_\epsilon U_\epsilon = P_r A P_c$ is the **computed** factorization of $P_r A P_c$ using the selected pivoting strategy. This is important in that it is used to assess the numerical stability of the chosen method to reliably computing the factorization and solving a system. This can be large for a poor choice of pivoting strategy even for well-conditioned matrices, i.e., ones that are not inherently difficult to factor accurately. This use will be discussed when we cover the numerical aspects of factorizations but it is also useful for checking the correctness of structured problems such as the one in the study questions with the large growth of elements.

Generation of Test Problems

A key consideration in this assignment is the generation of test problems. They must have full rank, i.e., linearly independent columns.

Some suggestions follow:

1. The matrices should be generated using double precision storage and computation. You can then coerce the type to single when storing the matrix before calling your single precision code if you decide to check both single and double precision execution of your codes. This is, however, not required but it is encouraged that you do this at some point.
2. Generate L and U so they are nonsingular unit lower triangular and upper triangular matrices respectively. Evaluate their product to define A . By their structure these matrices have linearly independent columns and therefore their product has full-rank. This is useful for both small and large values of n . For small values you can also constrain L and U to have integer values so A will have integer values. Take care with the magnitude of the elements of L and U . Nicely conditioned problems tend to be specified by off-diagonal elements in L smaller than 1 in magnitude and diagonal elements in U that have a reasonable range in magnitudes.¹ The off-diagonal elements in U should not be significantly larger in magnitude than the diagonal elements of U .
3. Remember even if A is generated from L and U when you run your routine with pivoting, nontrivial permutation matrices P_r and P_c may result and you may return \tilde{L} and \tilde{U} such that $P_r A P_c = \tilde{L} \tilde{U}$.
4. Square randomly generated matrices tend may be nonsingular and reasonably conditioned. Random triangular matrices tend not to be well-conditioned and solving them numerically may yield poor results.

¹Conditioning refers to a property of nonsingular matrices that indicates how sensitive the problem is to perturbations. This is used to alert a user to problems that may cause accuracy problems even for numerically robust and stable algorithms. You are encouraged to use the appropriate library routine, e.g., `cond(A)` in MATLAB, to compute the condition number of some of your test matrices. A condition number on the order of 10^k for moderate k is reasonable. A value of k around 6 or 7 and above may cause accuracy difficulties.

You can make sure any matrix is nonsingular by adding to the diagonal elements until the matrix is diagonally dominant by rows, columns or both.² This guarantees success of the factorization if it is run without pivoting. As noted above, if you allow pivoting the routine may so do even for a diagonally dominant matrix depending on the form of dominance and the pivoting strategy used. It is also useful to note that after generating such a matrix, A you can apply random permutations \tilde{P} and \hat{P} to generate a new test matrix $\tilde{A} = \tilde{P}A\hat{P}$ that will not be diagonally dominant but will still be nonsingular. Of course, \tilde{P} and \hat{P} are not necessarily the permutations that will be generated by applying partial or complete pivoting to \tilde{A} .

5. You can generate a symmetric positive definite $A \in \mathbb{R}^{n \times n}$ from a lower triangular \tilde{L} with positive diagonal elements (not necessarily 1) via $A = \tilde{L}\tilde{L}^T$. A is nonsingular by definition and factorization will succeed without pivoting. Note that your code will still produce an LU factorization since your factorization routine is not designed to exploit symmetry. However, there is a relationship between L , U and \tilde{L} . This is probed in the first set of structured factorization tasks.
6. Matrices with known structures that influence the structure or magnitude pattern of their factorizations are also useful. This is especially true if the patterns scale in a known way with n . Recall the example in the homework problems that has large elements in U . Consider what should happen when no pivoting or partial pivoting is used for various n values. Also, nonzero patterns such as banded matrices for A should produce specific nonzero patterns in L and U . Structure is the subject of the first set of empirical tasks discussed below.
7. Make sure that the matrices you use to check pivoting **actually require some pivoting when factored**.
8. You should run a range of problem sizes for each algorithm and problem type you evaluate.
9. Do not simply report the the factors or accuracy of the factorization for a small number of small systems. Think about how you would report the results of testing each of the routines with many matrices including those of sizes too large to display for any useful effect.

²A matrix is strictly diagonally dominant by rows (columns) if the magnitude of each diagonal element is strictly larger than the sum of the magnitudes of all off-diagonal elements in the same row (column). A matrix may of course be diagonally dominant by rows and columns simultaneously.

Empirical Tasks Set 1 : Structure

The first task involves matrices and factorizations with structure that can be deduced and checked. These are useful as initial debugging and validation examples.

1. Consider a matrix $A \in \mathbb{R}^{n \times n}$ that is diagonal with positive elements, i.e., $\alpha_{ij} = e_i^T A e_j = 0$ for all $i \neq j$ and $\alpha_{ii} > 0$ for $1 \leq i, j \leq n$. For example, let $\alpha_{ii} = i$ or $\alpha_{ii} = n - i + 1$

$$n = 5 \rightarrow A = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 4 & 0 \\ 0 & 0 & 0 & 0 & 5 \end{pmatrix} \quad \text{and} \quad A = \begin{pmatrix} 5 & 0 & 0 & 0 & 0 \\ 0 & 4 & 0 & 0 & 0 \\ 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 2 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{pmatrix}.$$

Consider factoring with no pivoting, partial row pivoting and complete pivoting. What can be said about the L , U , P_r and P_c factors as a function of your choice of structure in the positive diagonal elements? What is the growth factor for your chosen problems and the pivoting choices?

2. Consider a matrix $A \in \mathbb{R}^{n \times n}$ that is antidiagonal with positive elements, i.e., $\alpha_{1,n} > 0, \alpha_{2,n-1} > 0, \dots, \alpha_{n-1,2} > 0, \alpha_{n,1} > 0$. For example,

$$n = 5 \rightarrow A = \begin{pmatrix} 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 2 & 0 \\ 0 & 0 & 3 & 0 & 0 \\ 0 & 4 & 0 & 0 & 0 \\ 5 & 0 & 0 & 0 & 0 \end{pmatrix} \quad \text{and} \quad A = \begin{pmatrix} 0 & 0 & 0 & 0 & 5 \\ 0 & 0 & 0 & 4 & 0 \\ 0 & 0 & 3 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \end{pmatrix}.$$

Consider factoring with no pivoting, partial row pivoting and complete pivoting if you choose to implement it. What can be said about the L , U , P_r and P_c factors as a function of your choice of structure in the positive antidiagonal elements? What is the growth factor for your chosen problems and pivoting choices?

3. Consider a matrix $A \in \mathbb{R}^{n \times n}$ that is the sum of a diagonal matrix and an antidiagonal matrix with positive elements, i.e., and X nonzero pattern. Consider factoring with no pivoting, partial row pivoting and complete pivoting. What can be said about the L , U , P_r and P_c factors as a function of your choice of structure in the positive antidiagonal elements? If complete pivoting is there any structure that follows for a particular choice of pivot elements?
4. Consider a matrix $A \in \mathbb{R}^{n \times n}$ that is unit lower triangular and the elements in the strict lower part all have magnitude less than 1, i.e., $|\lambda_{ij}| < 1$ for $i > j$, $\lambda_{ij} = 0$ for $i < j$, and $\lambda_{ii} = 1$. Consider factoring with no pivoting, partial row pivoting and complete pivoting. What can be said about the L , U , P_r and P_c factors? (Note that your code does not know that A is unit lower triangular and will eliminate the elements below the diagonal by applying Gauss transforms.)

5. Consider a lower triangular matrix A again but this time let the diagonal elements be positive and not 1 and the elements in the strictly lower triangular part be larger than 1, e.g.,

$$A = \begin{pmatrix} 2 & 0 & 0 & 0 & 0 \\ 3 & 2 & 0 & 0 & 0 \\ 4 & 3 & 2 & 0 & 0 \\ 5 & 4 & 3 & 2 & 0 \\ 6 & 5 & 4 & 3 & 2 \end{pmatrix}.$$

Consider factoring with no pivoting, partial row pivoting and complete pivoting. What can be said about the L , U , P_r and P_c factors?

6. Consider a matrix $A \in \mathbb{R}^{n \times n}$ that is tridiagonal with all elements on the main diagonal and the first super and subdiagonals nonzero, i.e., $\alpha_{ii} \neq 0$, $\alpha_{i+1,i} \neq 0$, and $\alpha_{i,i+1} \neq 0$. Of course, these must be such that the matrix is nonsingular.

Suppose A is, additionally, strictly diagonally dominant by rows and columns. Consider factoring with no pivoting, partial row pivoting and complete pivoting. What can be said about the L , U , P_r and P_c factors?

7. Consider a matrix $A \in \mathbb{R}^{n \times n}$ with

- $\alpha_{ij} = e_i^T A e_j = -1$ when $i > j$, i.e., all elements strictly below the diagonal are -1 ;
- $\alpha_{ii} = e_i^T A e_i = 1$, i.e., all elements on the diagonal are 1;
- $\alpha_{in} = e_i^T A e_n = 1$, i.e., all elements in the last column of the matrix are 1;
- all other elements are 0

e.g., for $n = 4$ we have

$$A = \begin{pmatrix} 1 & 0 & 0 & 1 \\ -1 & 1 & 0 & 1 \\ -1 & -1 & 1 & 1 \\ -1 & -1 & -1 & 1 \end{pmatrix}$$

Consider factoring with no pivoting, partial row pivoting and complete pivoting. What can be said about the L , U , P_r and P_c factors? What is the growth factor for the different pivoting choices?

8. Suppose $A \in \mathbb{R}^{n \times n}$ is a symmetric positive definite generated from a lower triangular \tilde{L} with positive diagonal elements (not necessarily 1) via $A = \tilde{L}\tilde{L}^T$. A is nonsingular by definition and factorization will succeed without pivoting. Note that your code will still produce an LU factorization since your factorization routine is not designed to exploit symmetry. What is the relationship between L , U and \tilde{L} ?

Empirical Tasks Set 2 : General Trends

As with the assignment for Program 1, this set requires the generation of a large set of problems grouped and empirically analyzed by problem size n and, if appropriate the class of matrix problem considered.

For each problem size and class of problem, generate many example problems and evaluate the various metrics discussed earlier. You should present your results in a form appropriate to characterize these metrics over a large data set, i.e., too large to look at each problem individually. This can be done, for example, by graphs and histograms. The latter is particularly useful for detecting outliers in the performance such as large factorization error. These outliers can be discussed in more detail and explained if you wish.